






ordinary fair training




robust ensemble









DRO [Wang et al, NeurIPS 2020] hTERM [Li et al, ICLR 2021] [Konstantinov et al, ICML 2020]



FLEA (proposed)



accuracy



fairness